

Ethics of Artificial Intelligence in Crime

Kavya Sriram
kasriram@ucsd.edu

Catherine Back
cback@ucsd.edu

Aj Falak
pfalak@ucsd.edu

Yuancheng Cao
yuc094@ucsd.edu

Mentor: Emily Ramond
eramond@deloitte.com

Mentor: Greg Thein
gthein@deloitte.com

Background

Predictive models in pretrial risk assessment influence judicial decisions but often inherit racial biases from historical criminal justice data. This work examines racial bias in these models and applies bias mitigation techniques to improve fairness.

Pretrial Risk-Assessment Algorithms

- Predicts a defendant's risk of failing to appear or reoffending, influencing bail and detention decisions.
- Aims to reduce subjectivity but often reinforces systemic biases.

Bias in Risk Assessments

- COMPAS analysis (ProPublica, 2016) found Black defendants were twice as likely as White defendants to be falsely labeled high-risk.
- Bias originates from policing practices and socioeconomic disparities embedded in arrest data.

Data

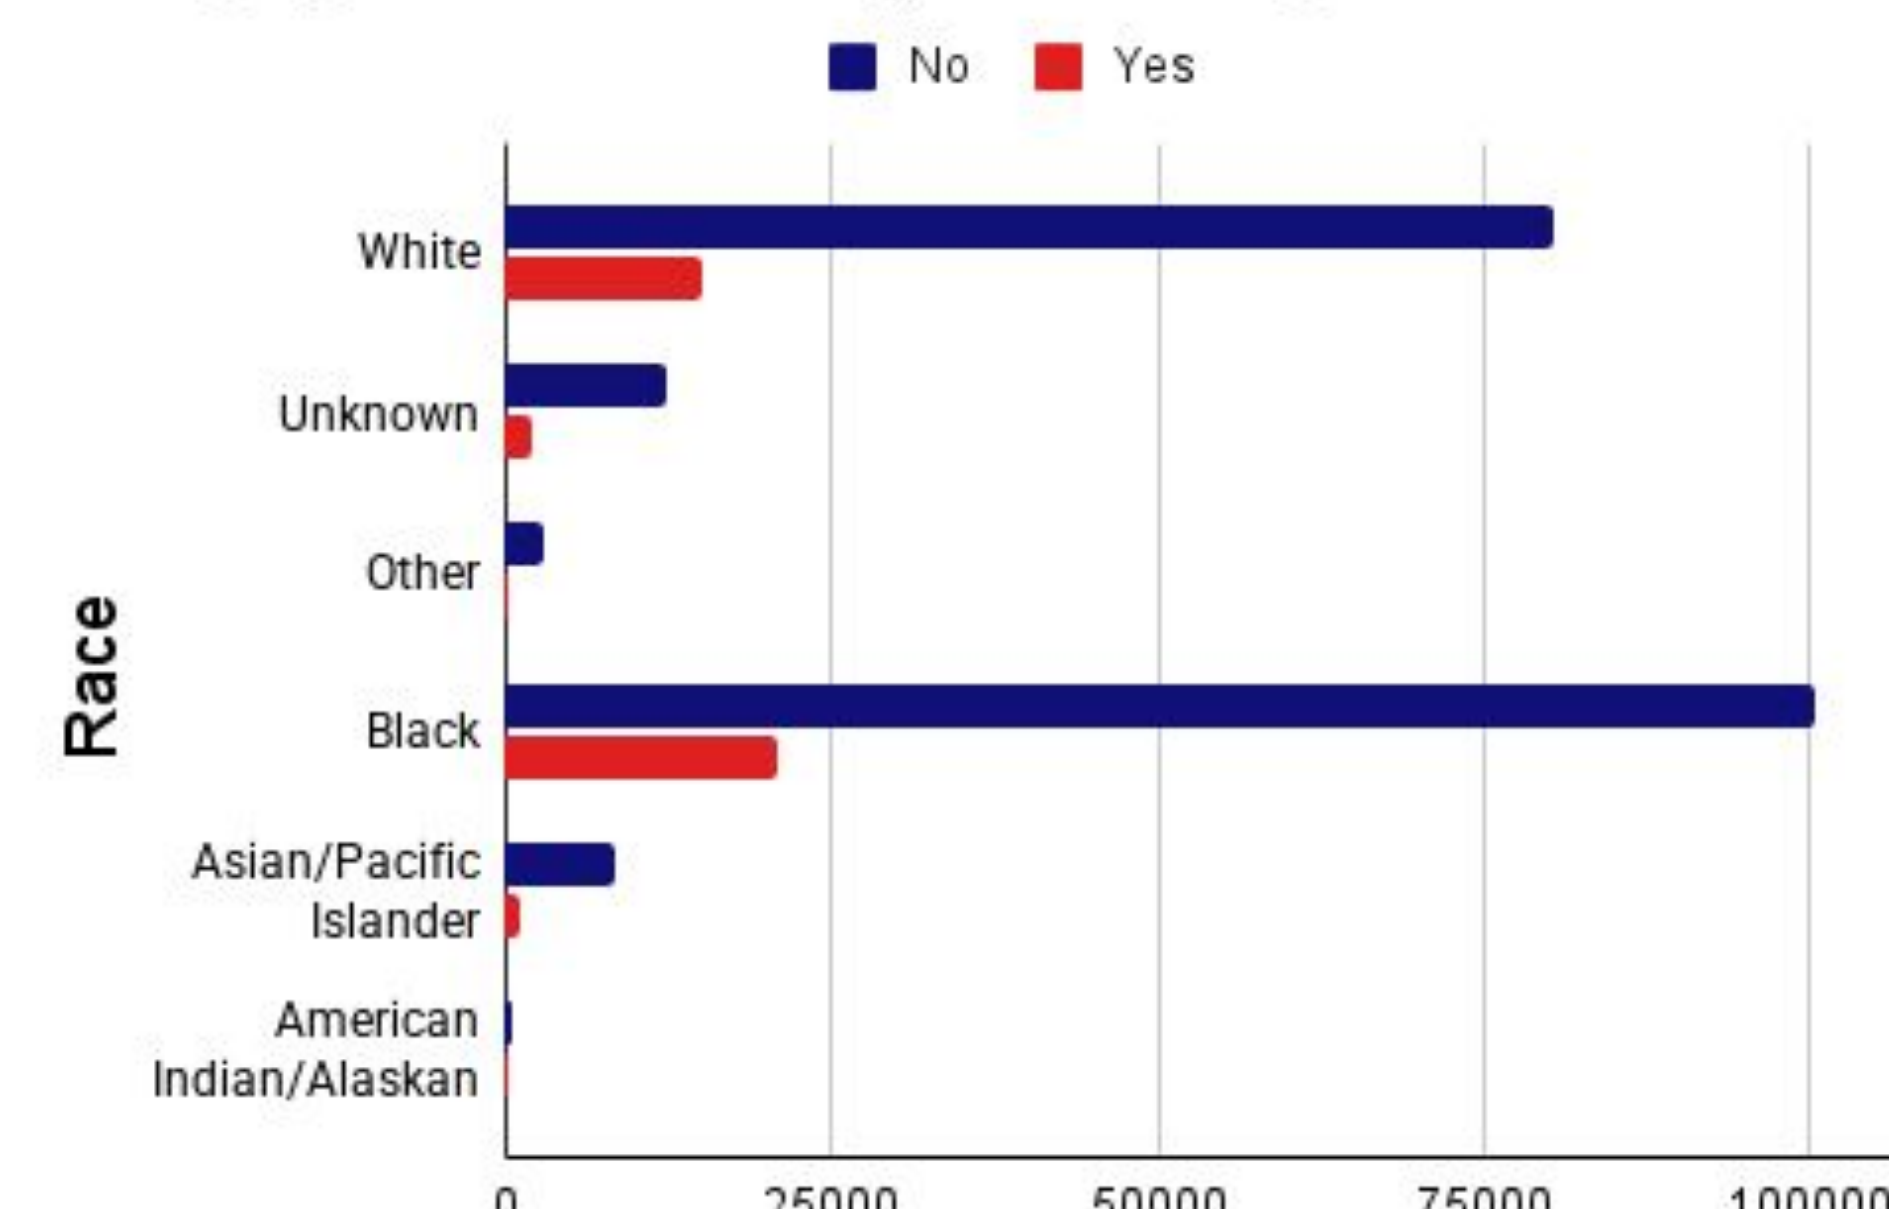
- Created with the Division of Criminal Justice Services (DCJS)
- Contains 244,271 records from 2023, including protected attributes like race.

Key Features

- Demographics: Race, ethnicity, gender, age at arrest/crime.
- Pretrial Decisions: Bail set/posted, release type, supervision type.
- Outcomes: Failure to Appear (FTA), reoffended, release decision, rearrest.
- Financial Factors: Bail amount, bond type.

Exploratory Data Analysis

Demographic Breakdown by Reoffending Status: Race



- The majority of individuals (83.7%) had no arrests during their pretrial period
- Black individuals accounted for 49.5% of cases, followed by White individuals at 38.9%, Hispanic population represented 24.8% of cases, while Non-Hispanic individuals accounted for 65.2%
- Black defendants showed a higher rate of pretrial rearrest (19.3%) compared to White defendants (14.7%) and Asian/Pacific Islander defendants (11.2%)

Model Development

- Built a Random Forest model using prior offences, pending charges, and crime severity, with binary indicators, 100 estimators, and a fixed random state for reproducibility
- Integrated SMOTE (Synthetic Minority Over-sampling Technique) with a Random Forest Classifier and setting class_weight='balanced' to adjust for class distribution
- Applied Stratified K-Fold Cross-Validation (k=5) and performed Grid Search to optimize parameters


Bias Mitigation

Reweighting (Pre-processing)

Assigns different weights to groups to correct imbalances to prevent bias toward the majority groups

Calibrated Equalized Odds (Post-processing)

Modified final model predictions to balance false positive and false negative rates across groups to ensure fair classification.



Results

- The original Random Forest Classifier achieved 83.3% accuracy but had a balanced accuracy of 50.7% indicating bias toward non-offenders
- After applying oversampling and fine-tuning, balanced accuracy improved to 64.66%

	Before Fine-Tuning	After Fine-Tuning
Recall (Class 1)	0.50	0.66
Recall (Class 0)	0.73	0.63
Precision (Class 1)	0.26	0.26
Precision (Class 0)	0.88	0.90

- Statistical analysis revealed racial bias in the model, with Black individuals having the highest predicted reoffender and false positive rates
- After Reweighting → Improved fairness, but a slight drop in accuracy due to balancing efforts.
- Postprocessing did not significantly impact balanced accuracy, meaning the model maintained its overall predictive performance.

Limitations & Next Steps

- Limited access to existing criminal justice models
- Enhance model accuracy while mitigating bias and expand to diverse datasets for broader applicability and fairness

Reference

Angwin, Julia, Jeff Larson, Surya Mattu, and Lauren Kirchner. 2016. "Machine Bias: There's Software Used Across the Country to Predict Future Criminals. And It's Biased Against Blacks." ProPublica.

Bellamy, Rachel K. E., et al. *AI Fairness 360: An Extensible Toolkit for Detecting, Understanding, and Mitigating Unwanted Algorithmic Bias*.



Learn More Here

